

Relational Algebra

COMP9311 24T2; Week 2.2

By Zhengyi Yang, UNSW

Motivation

- We've seen what a relational model is.
- We needed a formal language to specify data (tuples) from the relational model.
- **Relational Algebra** (E.F. Codd (1970))

Why Relational Algebra?

- It provides a formal foundation for relational model operations
- It is used as a basis for implementing and optimizing queries in the query processing and optimization
- Some of its concepts are incorporated into SQL

Relational Algebra

Relational Algebra is a procedural data manipulation language (**DML**).
It specifies operations on relations to define new relations:

Unary Relational Operations: Select, Project

Operations from Set Theory: Union, Intersection, Difference,
Cartesian Product

Binary Relational Operations: Join, Divide.

1 SELECT

The SELECT operation/predicate is used to select a subset of the tuples of a relation R, satisfying some conditions.

Notation: $\sigma_{\langle \text{selection condition} \rangle}(R)$

Intuition: Filters out all tuples that do not satisfy select condition



Selection Condition

The condition is defined by a ***selection clause***:

- *<attribute> operator <constant>*
- *< attribute > operator < attribute >*

Where *operator* is one of =, <, ≤, >, ≥ or ≠ ...

Example:

- *age ≤ 24*
- *commission ≥ 24 000*

Selection Condition

Selection clauses can also be

- `<expression> operator <expression>`

With this, we can use **Boolean connectives** as operators

- `C1 AND C2`
- `C1 OR C2`
- `NOT C`

Terms equivalently expressed by \wedge (and), \vee (or), \neg (not)

Q: Select the enrolment records for the students whose supervisor is Person 1

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

$\sigma_{(Supervisor=1)}(ENROLMENT)$

The output relation is

Enrolment#	Supervisee	Supervisor	Department	Degree
2	3	1	Comp.Sci	Ph.D.
3	4	1	Comp.Sci	M.Sc.
4	5	1	Comp.Sci	M.Sc.

Q: Select the enrolment records for Person 1's non-Ph.D. students

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

$\sigma_{(Supervisor=1 \text{ AND } Degree \neq "Ph.D.")}(ENROLMENT)$

$\sigma_{(Supervisor=1 \text{ AND } NOT Degree="Ph.D.")}(ENROLMENT)$

} Same

The output relation is

Enrolment#	Supervisee	Supervisor	Department	Degree
3	4	1	Comp.Sci	M.Sc.
4	5	1	Comp.Sci	M.Sc.

Properties of Selection

Properties:

- Consecutive selects ***can be combined***:

$$\sigma_{\langle cond1 \rangle}(\sigma_{\langle cond2 \rangle}(R)) = \sigma_{\langle cond1 \rangle \text{ AND } \langle cond2 \rangle}(R)$$

- Selection is a ***commutative*** operation:

$$\sigma_{\langle cond1 \rangle}(\sigma_{\langle cond2 \rangle}(R)) = \sigma_{\langle cond2 \rangle}(\sigma_{\langle cond1 \rangle}(R))$$

2 PROJECT

The PROJECT operation is used to project a subset of the attributes (column) of a relation, denoted by:

General form: $\pi_{\langle attribute\ list \rangle}(R)$

Result:

- schema: attribute list (A_1, \dots, A_k)
- instance: the set of all subtuples $t[A_1, \dots, A_k]$ where $t \in R$

Q: Find departments and degree requirements for the courses that students enroll.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

$\pi_{\{department, degree\}}(ENROLLMENT)$

The output relation is

Department	Degree
Psychology	Ph.D.
Comp.Sci	Ph.D.
Comp.Sci	M.Sc.

Duplicates of PROJECT

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

Question: What if we do PROJECTION on only department?

Department
Psychology
Comp.Sci.
Comp.Sci.
Comp.Sci.

or

Department
Psychology
Comp.Sci.

?

Duplicates of PROJECT

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

Question: What if we do PROJECTION on only department?

Answer: Keep only **one** 'Comp.Sci.'.

Department

Psychology

Comp.Sci.

Relational Algebra is based on sets, so no duplicates are allowed.

- The PROJECT operation *removes any duplicate tuples*, so the result of the PROJECT operation is a set of distinct tuples, and this is known as **duplicate elimination**.

Properties of PROJECT

Consider $\pi_{\langle list1 \rangle}(\pi_{\langle list2 \rangle}(R))$

If $\langle list2 \rangle$ contains all the attributes in $\langle list1 \rangle$:

Then $\pi_{\langle list1 \rangle}(\pi_{\langle list2 \rangle}(R)) = \pi_{\langle list1 \rangle}(R)$

Else the operation is *not well defined*.

Project Predicate

Question: is projection commutative with selection?

$$\text{i.e., } \pi_X(\sigma_B(R)) = \sigma_B(\pi_X(R))?$$

Consider the following:

$$\pi_{\{degree\}}(\sigma_{(Department='Psychology')}(ENROLMENT))$$

Degree
Ph.D.


$$\sigma_{(Department='Psychology')}(\pi_{\{degree\}}(ENROLMENT))$$

Error as SELECT cannot find Department

Answer: The attribute used in SELECT must be a subset of the attribute list in PROJECT

Intuition: Projection and Selection

1. Selection performs a horizontal decomposition, and
2. projection performs a vertical decomposition



3 SET UNION

UNION is the set-theoretic union of the tuples of two relations.

$$R \cup S = \{t: t \in R \text{ or } t \in S\}$$

Condition: R and S must be **union compatible**!

Union compatibility: there is a 1-1 correspondence between their attributes: the same name and same domain.

Example: to find all courses taught in the Fall 2009 semester, **or** in the Spring 2010 semester, **or** in both:

$$\pi_{\{course_id\}}(\sigma_{(semester="Fall" \wedge year=2009)}(section)) \cup \pi_{\{course_id\}}(\sigma_{(semester="Spring" \wedge year=2010)}(section))$$

Example

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Example: $STUDENT \cup RESEARCHER =$

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledhill
5	Ms B.K.Lee
2	Dr R.G.Wilkinson

4 SET INTERSECTION

- *INTERSECTION* is an operation that includes all tuples that are in present both relations, denoted by

$$R \cap S = \{t: t \in R \text{ and } t \in S\}$$

- Condition: R and S must also be **union compatible!**

- Example: $R_1 \leftarrow \sigma_{(supervisor=1)}(ENROLMENT)$
 $R_2 \leftarrow \sigma_{(degree='Ph.D.')} (ENROLMENT)$

$$R_1 \cap R_2 =$$

Enrolment#	Supervisee	Supervisor	Department	Degree
2	3	1	Comp.Sci.	Ph.D.

Example of Intersection

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Example: $\text{STUDENT} \cap \text{RESEARCHER} =$

Person#	Name
1	Dr C.C. Chen

5 SET DIFFERENCE

DIFFERENCE is a relation that includes all tuples that are in the left relation but not in the right relation, denoted by

$$R - S = \{t: t \in R \text{ and } t \notin S\}$$

Condition: R and S must also be **union compatible**!

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson


Example: STUDENT – RESEARCHER =

Person#	Name
3	Ms K. Juliff
4	Ms J. Gledhill
5	Ms B.K. Lee


Summary

Operations on Relations


- *UNION:* $A \cup B$



- *INTERSECTION:* $A \cap B$



- *DIFFERENCE:* $A - B$



Express: The names of persons who are either a student **or** a researcher

STUDENT:

<u>Person#</u>	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Express: The names of persons who are either a student **or** a researcher

STUDENT:

<u>Person#</u>	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

$\pi_{\{name\}}(STUDENT \cup RESEARCHER)$

Name
Dr C.C.Chen
Dr R.G.Wilkinson
Ms K.Juliff
Ms J.Gledill
Ms B.K.Lee

Express: The names of persons who are a student and a researcher

STUDENT:

<u>Person#</u>	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Express: The names of persons who are a student and a researcher

STUDENT:

<u>Person#</u>	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

$\pi_{\{name\}}(STUDENT \cap RESEARCHER)$

Name
Dr C.C.Chen

Express: The names of persons who are a student **but not** a researcher

STUDENT:

<u>Person#</u>	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

Express: The names of persons who are a student **but not** a researcher

STUDENT:

<u>Person#</u>	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

$\pi_{\{name\}}(STUDENT - RESEARCHER)$

Name
Ms K.Juliff
Ms J.Gledill
Ms B.K.Lee

Express: The departments and degrees of Courses which are not enrolled by any student

ENROLMENT:

<u>Enrolment#</u>	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

COURSE:

<u>Department</u>	<u>Degree</u>
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

Express: The departments and degrees of Courses which are not enrolled by any student

ENROLMENT:

<u>Enrolment#</u>	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

COURSE:

<u>Department</u>	<u>Degree</u>
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

$COURSE - (\pi_{(Department, degree)}(ENROLMENT))$

CARTESIAN PRODUCT

$$R \times S = \{t_1 || t_2 : t_1 \in R \text{ and } t_2 \in S\}$$

- Intuition: **every combination of tuples in R with tuples in S.**
- $t_1 || t_2$ indicates the concatenation of tuples.
- R and S not required to be union compatible, but
- The number of tuples in the output relations is always $|R| * |S|$

Usually assumes that attributes of $r(R)$ and $s(S)$ are disjoint. (That is, $R \cap S = \emptyset$). If not, you must devise a naming schema to distinguish between the attribute names if they are the same in $r(A, B)$ and $s(A, C)$, by attaching the relation's name, $r.A$ and $s.A$ (known as **dot-notation**)

Example of cartesian product

ENROLMENT:

<u>Enrolment#</u>	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

RESEARCHER:

<u>Person #</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson



Example of cartesian product

ENROLMENT X RESEARCHER =

<u>E'ment#</u>	S'ee	S'or	D'ment	Degree	Person #	Name
1	1	2	Psych.	Ph.D.	1	Dr C.C. Chen
1	1	2	Psych.	Ph.D.	2	Dr R.G.Wilkinson
2	3	1	Cmp.Sci	Ph.D.	1	Dr C.C. Chen
2	3	1	Cmp.Sci	Ph.D.	2	Dr R.G.Wilkinson
3	4	1	Cmp.Sci	M.Sc.	1	Dr C.C. Chen
3	4	1	Cmp.Sci	M.Sc.	2	Dr R.G.Wilkinson
4	5	1	Cmp.Sci	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci	M.Sc.	2	Dr R.G.Wilkinson

There were 4 tuples in ENROLMENT and 2 tuples in RESEARCHER. In the result, there are 8 tuples.

Useful if we add a condition

$$R_1 \leftarrow ENROLMENT \times RESEARCHER$$

<u>E'ment#</u>	S'ee	S'or	D'ment	Degree	Person#	Name
1	1	2	Psych.	Ph.D.	1	Dr C.C. Chen
1	1	2	Psych.	Ph.D.	2	Dr R.G. Wilkinson
2	3	1	Cmp.Sci	Ph.D.	1	Dr C.C. Chen
2	3	1	Cmp.Sci	Ph.D.	2	Dr R.G. Wilkinson
3	4	1	Cmp.Sci	M.Sc.	1	Dr C.C. Chen
3	4	1	Cmp.Sci	M.Sc.	2	Dr R.G. Wilkinson
4	5	1	Cmp.Sci	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci	M.Sc.	2	Dr R.G. Wilkinson

In practice it's useful if we give a cartesian product specified condition

$$\sigma_{(Supervisor=Person\#)}(R_1) =$$

E'ment#	S'ee	S'or	D'ment	Degree	Person#	R'cher. Name
1	1	2	Psych.	Ph.D.	2	Dr R.G. Wilkinson
2	3	1	Cmp.Sci.	Ph.D.	1	Dr C.C. Chen
3	4	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen

More useful if we add a projection

$R_1 \leftarrow ENROLMENT \times RESEARCHER$

$R_2 \leftarrow \sigma_{(Supervisor=Person\#)}(R_1)$

E'ment#	S'ee	S'or	D'ment	Degree	Person#	R'cher. Name
1	1	2	Psych.	Ph.D.	2	Dr R.G.Wilkinson
2	3	1	Cmp.Sci.	Ph.D.	1	Dr C.C. Chen
3	4	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen
4	5	1	Cmp.Sci.	M.Sc.	1	Dr C.C. Chen

$\pi_{\{E'ment\#,S'ee,S'or,Name,D'ment,Degree\}}(R_2)$

E'ment#	S'ee	S'or	Name	D'ment	Degree
1	1	2	Dr R.G.Wilkinson	Psych.	Ph.D.
2	3	1	Dr C.C. Chen	Comp.Sci.	Ph.D.
3	4	1	Dr C.C. Chen	Comp.Sci.	M.Sc.
4	5	1	Dr C.C. Chen	Comp.Sci.	M.Sc.

The two equal attributes occur only once

The last of these is also known as *natural join*, the next to last is *equi-join*.

6 JOIN

- JOIN is used to combine related tuples from two relations into single "longer" tuples.
- **Theta-join**

$$R \bowtie_{\langle \text{join condition} \rangle} S = \{t_1 \parallel t_2 : t_1 \in R \text{ and } t_2 \in S \text{ and } \langle \text{join condition} \rangle\}$$

- A general join condition is of the form:

$\langle \text{condition} \rangle$ **AND** $\langle \text{condition} \rangle$ **AND** ... **AND** $\langle \text{condition} \rangle$

6.1 Equi-join

A type of theta-join where the only comparison operator used is “=” is called an Equi-join

Example:

$$ENROLMENT \bowtie_{(Supervisor=Person\#)} RESEARCHER$$

6.2 Natural Join

A type of equi-join that requires each pair of join attributes to have the same name and domain in both relations.

Notes: In a natural join, there may be several valid pairs of join attributes.

$ENROLMENT \bowtie_{(department, name), (deparment, name)} COURSE$

If there are pairs of joining attributes identically named, we can write

$ENROLMENT \bowtie COURSE$

Note: this notion also acceptable if there's one join attribute

6.2 Natural Join

Intuitions:

- Enforce equality on all attributes with same name
- Eliminate one copy of duplicated attributes

JOINS

Remember the differences between the types of joins:

1. Theta JOIN
2. Equi JOIN
3. Natural JOIN

Note: all denoted with \bowtie

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

What are the names of students who are studying for an MSc in computer science?

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

What are the names of students who are studying for an MSc in computer science?

$\pi_{\{name\}}(\sigma_{(degree=MSc \text{ and } Depart=CS)} ENROLMENT \bowtie_{supervisee=person\#} Student)$

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

The IDs of students who are supervised by Dr C.C.Chen

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

The IDs of students who are supervised by Dr C.C.Chen

$R1 = \text{ENROLMENT} \bowtie (\text{supervisor}=\text{person\#}) \text{ RESEARCHER}$

$R2 = \sigma_{(\text{name}=\text{Dr C.C.Chen})} R1$

$R3 = \pi_{\{\text{supervisee}\}} R2$

Divide

- The DIVISION operation is applied to two Relations R and S, where the attributes of S are a subset of the attributes of R.
- The relation returned by the division operator will have attributes = (All attributes of R – All Attributes of S)
- Return all tuples from relation R which are associated to every S's tuple.

R	A	B	S
	a ₁	b ₁	
	a ₁	b ₂	
	a ₂	b ₁	
	a ₃	b ₂	
	a ₄	b ₁	
	a ₅	b ₁	
	a ₅	b ₂	
$R \div S =$			
A			
a ₁			
a ₅			

Divide

Typical use: which courses are offered by all departments?

$$Course \div (\pi_{Department} Course)$$

Divide

Typical use: which courses are offered by all degrees?

$Course \div (\pi_{Degree} Course)$

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

The names of supervisor who supervises both MSc and PhD students

Pracs

STUDENT:

<u>Person#</u>	Name
1	Mr J.He
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

<u>Person#</u>	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE

<u>Depart</u>	<u>Degree</u>
EE	PhD
CS	PhD
EE	MSc
CS	MSc

ENROLMENT:

<u>Enrol#</u>	Supervisee	Supervisor	Depart	Degree
1	1	2	EE	PhD
2	3	1	CS	PhD
3	4	1	CS	MSc
4	5	1	CS	MSc

The names of supervisor who supervises both MSc and PhD students

$R1 = \pi_{\{SUPERVISOR, DEGREE\}} ENROLMET \div \pi_{\{DEGREE\}} COURSE$

$R2 = \pi_{\{Name\}} (R1 \bowtie_{(supervisor=person\#)} RESEARCH)$

Exercise

R:

A	B	C
a ₁	b ₁	c ₁
a ₁	b ₁	c ₂
a ₁	b ₁	c ₃
a ₁	b ₂	c ₂
a ₂	b ₁	c ₁
a ₂	b ₂	c ₂
a ₃	b ₁	c ₁
a ₃	b ₂	c ₁
a ₃	b ₂	c ₂

S:

B	C
b ₁	c ₁
b ₁	c ₂

Write relational algebra that retrieves:

1. Find A of R that contains all S.
2. Find (A, B) of R that contains all C of S.

Exercise Answers:

1. $R \div S$

A
a ₁

2. $R \div \pi_{\{c\}}(S)$

A	B
a ₁	b ₁
a ₃	b ₂

Rename Operator

- The **rename** operator ρ changes the name of one or more attributes
- Change the names in a schema
- Does not affect **instance** of the target relation

Family

Father	Child
Adam	Abel
Adam	Cain
Abraham	Isaac

ρ (Parent, Child) (Family)

Parent	Child
Adam	Abel
Adam	Cain
Abraham	Isaac

- Why might this be useful? To be included in relational algebra?

Why RENAME Operator?

- To unify schemas for set operators
- For disambiguation in “self-join”

Basic vs Extended Operators

Note: $\{\sigma, \pi, \cup, -, \times\}$ (and *rename*) are sufficient to define all these operations: this is a relationally complete set of operators. These are the **basic operators** of the Relational Algebra.

What about *JOIN*, *INTERSECTION* and *DIVIDE*?

They are **extended operators** because they can be derived from the basic operators.

E.g., We can write $R \div S$ as

$$TEMP1 \leftarrow \pi_{R-S}(R)$$
$$TEMP2 \leftarrow \pi_{R-S}((TEMP1 \times S) - R)$$
$$RESULT = TEMP1 - TEMP2$$

- The result to the right of \leftarrow is assigned to the relation variable on the left of \leftarrow .
- May use variable in subsequent expressions.

Aggregate Operators

What if we want a relation with information about “sum of salaries” of employees, or the “average age” of students?

We need more expressive power, we can use **aggregation functions** to summarize information from multiple tuples into **aggregate values**.

We can use an **aggregation operator** γ and a function such as *SUM*, *AVG*, *MIN*, *MAX*, or *COUNT*. What if NULL?

If $R =$	A	B	, then $\gamma_{SUM(A)}(R) =$	SUM(A)
	1	2		8
	3	4		
	3	5		
	1	1		
			and $\gamma_{AVG(B)}(R) =$	AVG(B)
				3

Aggregate Operators

We can also retrieve aggregate values for groups, like the “sum of employee salaries” *per department* or the “average student age” *per faculty*.

We give γ additional arguments to specify that the aggregation behavior should be based on groups (defined by a set of attributes).

If $R =$

a	b
1	2
3	4
3	5
1	3

, then $\gamma_{a, \text{SUM}(b)}(R) =$

a	SUM(b)
1	5
3	9

Formal Definition

A **basic relational algebra expression** is one of the following:

- A relation in the database
- (could also be a) constant relation

A **general relational algebra expression** is constructed out of smaller subexpressions. Let E_1 and E_2 be relational algebra expressions; the following are all relational-algebra expressions:

- $E_1 \cup E_2$
- $E_1 - E_2$
- $E_1 \times E_2$
- $\sigma_P(E_1)$ where P is predicate on attributes in E_1
- $\pi_S(E_1)$ where S is a set of attributes in E_1
- $\rho_X(E_1)$ where X is the new name for the result of E_1

OPERATION	PURPOSE	NOTATION
SELECT	Selects all tuples that satisfy the selection condition from a relation R	$\sigma_{\langle \text{selection condition} \rangle}(R)$
PROJECT	Produces a new relation with only some of the attributes of R and removes duplicate tuples.	$\pi_{\langle \text{attribute list} \rangle}(R)$
THETA-JOIN	Produces all combinations of tuples from R and S that satisfy the join condition.	$R \bowtie_{\langle \text{join condition} \rangle} S$
EQUI-JOIN	Produces all the combinations of tuples from R and S that satisfy a join condition with only equality comparisons.	$R \bowtie_{\langle \text{join condition} \rangle} S$
NATURAL-JOIN	Same as EQUIJOIN except that the join attributes of S are not included in the resulting relation; if the join attributes have the same names, they do not have to be specified at all.	$R \bowtie_{\langle \text{join condition} \rangle} S$
UNION	Produces a relation that includes all the tuples in R or S or both R and S; R and S must be union compatible.	$R \cup S$
INTERSECTION	Produces a relation that includes all the tuples in both R and S; R and S must be union compatible.	$R \cap S$
DIFFERENCE	Produces a relation that includes all the tuples in R that are not in S; R and S must be union compatible.	$R - S$
CARTESIAN PRODUCT	Produces a relation that has the attributes of R and S and includes as tuples all possible combinations of tuples from R and S.	$R \times S$
DIVISION	Produces a relation T(X) that includes all tuples t[X] in R(Z) that appear in R in combination with every tuple from S(Y), where $Z = X \cup Y$.	$R(Z) \div S(Y)$