

# Revision

COMP9311 24T2; Week 10

*By Zhengyi Yang, UNSW*

# Details of Final Exam:

1. **Time**: 1:55PM to 5:10PM, Sydney time. 15th August 2024.
2. **Exam paper**: will be released on our course website & Moodle just before 1:55PM **15<sup>th</sup> Aug**.
3. **How to submit**: via **Moodle**
  - Same procedure as what you did for the two assignments.
  - In case near 5:00PM and you cannot submit via Moodle, send your solution to [junhua.zhang@unsw.edu.au](mailto:junhua.zhang@unsw.edu.au) and [zhengyi.yang@unsw.edu.au](mailto:zhengyi.yang@unsw.edu.au) with you zID and name.
  - You can submit multiple times and we will mark the last one.
  - We accept *any* format: directly answer using Word or handwriting and convert to Word/pdf. As long as the file is in .doc or .pdf format and **clear**.
4. Any question during the exam, send email to [junhua.zhang@unsw.edu.au](mailto:junhua.zhang@unsw.edu.au) and [zhengyi.yang@unsw.edu.au](mailto:zhengyi.yang@unsw.edu.au)

# Final Exam

- Time allowed:
  - 3 hrs + 15 mins for downloading/uploading (the additional 15 mins are not used to answer questions)
- **If you do not feel well on the exam day, please don't attend the exam.** By sitting or submitting an exam or assessment on the scheduled assessment date, you are declaring that you are fit to do so and cannot later apply for Special Consideration. (I.e., no sup-exam will be given.)

# Late Submission

You **MUST** submit your answers on time

**5% late penalty per minute** for late submissions.

**Submitting wrong files will result in 0 marks.**

# My Experience Survey

The UNSW MyExperience survey for Term 2 2024 is still open, and with a response rate of less than 15%, your participation is highly encouraged: <https://myexperience.unsw.edu.au/>

**Bonus** time will be announced.

# Plagiarism



★ We adopt a **zero-tolerance** policy for plagiarism.

All submissions are checked for plagiarism. The university regards plagiarism as a form of academic misconduct and has very strict rules regarding plagiarism.

For UNSW policies, penalties, and information to help avoid plagiarism, please see: <https://student.unsw.edu.au/plagiarism>. *Not knowing the rules is not considered a valid excuse.*

**All assessments must be your own original work. They are NOT group project.**

DO NOT: copy from others, copy from the Internet, pay someone to do it.

**You are not allowed to use generative AI tools such as ChatGPT.**

# Question Types:

- 6 questions
- Similar to questions in assignments and sample exam paper
- Sample exam paper is released on WebCMS

# Consultation:

We will run daily consultations (in-person and online) before the final exam:

- *Monday 5 Aug, 2pm-4pm, k17 203*
- *Tuesday 6 Aug, 2pm-4pm, k17 203*
- *Wednesday 7 Aug, 2pm-4pm, k17 203*
- *Thursday 8 Aug, 1pm-3pm, k17 203*
- *Friday 9 Aug, 2pm-4pm, k17 201B Meeting Room*
  
- *Monday 12 Aug, 1pm-3pm, k17 201B Meeting Room*
- *Tuesday 13 Aug, 4pm-6pm, k17 201B Meeting Room*
- *Wednesday 14 Aug, 12pm-2pm, k17 201B Meeting Room*



# Note 1:

## Computer Updates

You must ensure that auto-updates are **disabled** on your computer prior to the online assessment.

Special consideration will **NOT** be awarded on the grounds that your computer performed an update during an online assessment.

## Note 2:

**If you accidentally upload the wrong document or wrong version of your exam**

Students are responsible for uploading the **correct** version of the correct document. Once uploaded, there will be no opportunity to replace or re-upload your exam papers **AFTER** the end of the exam.

The documents submitted will be the documents that are marked. There is **NO** provision for students who upload incorrect or incomplete documents.

Therefore, you must check the work before you submit.

Check if you have submitted the correct file.

# Note 3:

## Communication during the exam

Students are **NOT** permitted to communicate with other people during the exam (including the reading and submission periods).

Attempts to communicate with other students will be considered to be serious academic misconduct.

This includes communication in person, by email, text, message, telephone, or internet.

i.e., do the work yourself

## Note 4:

### **Sharing answers with others or posting them online**

Any attempts to collaborate or share your answers with others will be considered a very serious case of academic misconduct

## Note 5 (Checklist):

1. Be logged in at your computer and ready to go 15 minutes before the exam commences.
2. Ensure your device has power, and the charger is plugged in.
3. If applicable remind your roommates or family that you'll be taking an exam to avoid interruptions.

## Note 6:

If you experience other issues, take **timestamped** screenshots

Contact the Course Coordinator or Admin via email to advise you are experiencing a technical issue, as soon as possible.

# Tips

- You can attempt the questions in any order, **arrange your time wisely.**
- Read all questions carefully
- Be fully prepared before the exam: don't forget to eat, take a break and relax yourself, no need to panic

# Overview: Database Design 1

## Data Models:

- ER, Relational Data Model and their mapping

## Relational Algebra:

- Be able to use relational algebra to answer question.

## Database Languages:

- SQL (final exam: may need to write some simple SQL)
- PL/pgSQL (not in exam)

## Relational Database Design:

- Functional Dependency
- Normal Forms
- Design Algorithms for 3NF and BCNF



# Overview: Database Design 2

## Data Storage:

- Record Format
- Buffer Management

## Query Optimisation (not in exam):

- Index
- Query Plan/Join Order Selection

## Transaction Management:

- Concurrency Control
- Recovery

## NoSQL :

- NoSQL Concept
- Different Data Model: Key-Value, Document, Column-family, Graph

## Guest Lecture (not in exam)

# Data Models

ER, Relational Data Model and their mapping

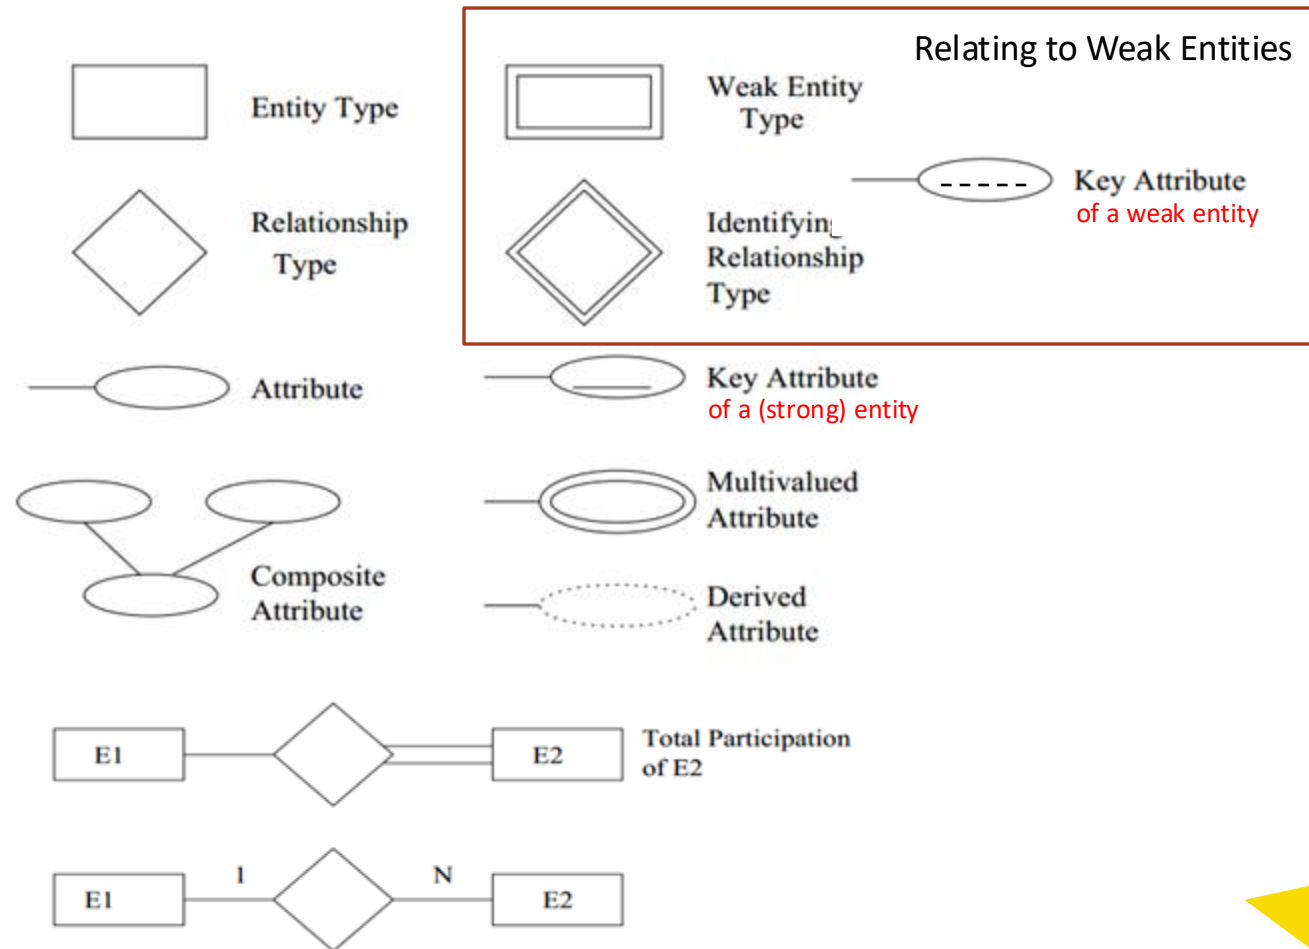


# Entity-Relationship Model

1. Entity Type: Group of object with the same properties
2. Entity: member of an entity type - analogous to an object.
3. Attribute: a property of object
4. Relationship: among objects

# Notations

The notation used for ERDs is summarised in Elmasre/Navathe Figure 3.15.



# Relational Data Model

- In the relational model, everything is described using relations.
- A relation can be thought of as a named table.
- Each column of the table corresponds to a named attribute.
- The set of allowed values for an attribute is called its domain.
- Each row of the table is called a tuple of the relation.
- N.B. There is no ordering of column or rows.

# Keys

- *Keys* are used to identify tuples in a relation.
- A *superkey* is a set of attributes that uniquely determines a tuple.
- Note that this is a property of the relation that does not depend on the current relation instance.
- A *candidate key* is a superkey, none of whose *proper* subsets is a superkey.
- Keys are determined by the applications.

# Integrity Constraints

There are several kinds of integrity constraints that are an integral part of the relational model:

1. *Key constraint*: candidate key values must be unique for every relation instance.
2. *Entity integrity*: an attribute that is part of a primary key cannot be NULL.
3. *Referential integrity*: The third kind has to do with “foreign keys”.

# Foreign Keys

Foreign keys are used to refer to a tuple in another relation.

A set, FK, of attributes from a relation schema R1 may be a foreign key if

- the attributes have the same domains as the attributes in the primary key of another relation schema R2, and
- a value of FK in a tuple t1 of R1 either occurs as a value of PK for some tuple t2 in R2 or is null.

Referential integrity: The value of FK must occur in the other relation or be entirely NULL.



# ER to Relational Model Mapping

One technique for database design is to first design a conceptual schema using a high-level data model, and then map it to a conceptual schema in the DBMS data model for the chosen DBMS.

Here we looked at a way to do this mapping from the ER to the relational data model. (see details in the lecture notes of Relational Data Model).

Composite and multivalued attributes are allowed in ER model, but not allowed in relational data model

# Relational Algebra

*Relational Algebra* is a procedural data manipulation language (**DML**). It specifies operations on relations to define new relations:

**Unary Relational Operations:** Select, Project

**Operations from Set Theory:** Union, Intersection, Difference,  
Cartesian Product

**Binary Relational Operations:** Join, Divide.

**Relational Algebra:** be able to use relational algebra to answer question.

OPERATION	PURPOSE	NOTATION
SELECT	Selects all tuples that satisfy the selection condition from a relation R	$\sigma_{\langle selection\ condition \rangle}(R)$
PROJECT	Produces a new relation with only some of the attributes of R and removes duplicate tuples.	$\pi_{\langle attribute\ list \rangle}(R)$
THETA-JOIN	Produces all combinations of tuples from R and S that satisfy the join condition.	$R \bowtie_{\langle join\ condition \rangle} S$
EQUI-JOIN	Produces all the combinations of tuples from R and S that satisfy a join condition with only equality comparisons.	$R \bowtie_{\langle join\ condition \rangle} S$
NATURAL-JOIN	Same as EQUIJOIN except that the join attributes of S are not included in the resulting relation; if the join attributes have the same names, they do not have to be specified at all.	$R \bowtie_{\langle join\ condition \rangle} S$
UNION	Produces a relation that includes all the tuples in R or S or both R and S; R and S must be union compatible.	$R \cup S$
INTERSECTION	Produces a relation that includes all the tuples in both R and S; R and S must be union compatible.	$R \cap S$
DIFFERENCE	Produces a relation that includes all the tuples in R that are not in S; R and S must be union compatible.	$R - S$
CARTESIAN PRODUCT	Produces a relation that has the attributes of R and S and includes as tuples all possible combinations of tuples from R and S.	$R \times S$
DIVISION	Produces a relation T(X) that includes all tuples t[X] in R(Z) that appear in R in combination with every tuple from S(Y), where $Z = X \cup Y$ .	$R(Z) \div S(Y)$

# Database Languages

- Database Languages: SQL
- **final exam: may need to write some simple SQL**

# SQL Queries<sub>(cont)</sub>

Query syntax is:

SELECT attributes

FROM relations

WHERE condition

The result of this statement is a table, which is typically displayed on output.

The SELECT statement contains the functionality of *select*, *project* and *join* from the relational algebra.

# SQL Comparisons

Comparison operators are defined on all types.

<   >   <=   >=   =   !=

Boolean operators AND, OR, NOT are available within WHERE expressions to combine results of comparisons.

Comparison against NULL yields FALSE.

Can explicitly test for NULL using:

- ***attr* IS NULL**                      ***attr* IS NOT NULL**

Most data types also have type-specific operations available (e.g., arithmetic for numbers).

# Relational Database Design

## Relational Database Design:

- Functional Dependency
- Normal Forms
- Design Algorithms for 3NF and BCNF

# Functional Dependencies

A function  $f$  from  $S_1$  to  $S_2$  has the property

$$\text{if } x, y \in S_1 \text{ and } x = y, \text{ then } f(x) = f(y).$$

A generalization of keys to avoid design flaws violating the above rule.

Let  $X$  and  $Y$  be sets of attributes in  $R$ .

$X$  (*functionally*) determines  $Y$ ,  $X \rightarrow Y$ , iff  $t_1[X] = t_2[X]$  implies  $t_1[Y] = t_2[Y]$ .

$$\text{i.e., } f(t(X)) = t[Y]$$

We also say  $X \rightarrow Y$  is a *functional* dependency, and that  $Y$  is *functionally* dependent on  $X$ .

$X$  is called the *left side*,  $Y$  the *right side* of the dependency.



# Armstrong's Axioms

*Notation:* If  $X$  and  $Y$  are sets of attributes, we write  $XY$  for their union.

e.g.,  $X = \{A, B\}$ ,  $Y = \{B, C\}$ ,  $XY = \{A, B, C\}$

1. F1 (Reflexivity) If  $X \supseteq Y$  then  $X \rightarrow Y$ .
2. F2 (Augmentation)  $\{X \rightarrow Y\} \models XZ \rightarrow YZ$ .
3. F3 (Transitivity)  $\{X \rightarrow Y, Y \rightarrow Z\} \models X \rightarrow Z$ .
4. F4 (Additivity)  $\{X \rightarrow Y, X \rightarrow Z\} \models X \rightarrow YZ$ .
5. F5 (Projectivity)  $\{X \rightarrow YZ\} \models X \rightarrow Y$ .
6. F6 (Pseudo-transitivity)  $\{X \rightarrow Y, YZ \rightarrow W\} \models XZ \rightarrow W$ .

# The Procedure for Computing $F^+$

To compute the closure of a set of functional dependencies  $F$ :

$F^+ = F$

**repeat**

**for each** functional dependency  $f$  in  $F^+$

        apply reflexivity and augmentation rules on  $f$

        add the resulting functional dependencies to  $F^+$

**for each** pair of functional dependencies  $f_1$  and  $f_2$  in  $F^+$

**if**  $f_1$  and  $f_2$  can be combined using transitivity

**then** add the resulting functional dependency to  $F^+$

**until**  $F^+$  does not change any further

# Algorithm to Compute $X^+$

An **algorithm** for you to follow step by step

```
X+ := X;  
change := true;  
while change do  
  begin  
    change := false;  
    for each FD  $W \rightarrow Z$  in  $F$  do  
      begin  
        if  $(W \subseteq X^+) \text{ and } (Z - X^+ \neq \emptyset)$  then do  
          begin  
             $X^+ := X^+ \cup Z$ ;  
            change := true;  
          end  
        end  
      end  
    end  
  end  
end
```

# Compute a Candidate Key

Given a relational schema  $R$  and a set  $F$  of functional dependencies on  $R$ .

A key  $X$  of  $R$  must have the property that  $X^+ = R$ .

## Algorithm

- Step 1: Assign  $X$  a superkey in  $F$ .
- Step 2: Iteratively remove attributes from  $X$  while retaining the property  $X^+ = R$ .
- The remaining  $X$  is a key.

# Compute All the Candidate Keys

Given a relational schema  $R$  and a set  $F$  of functional dependencies on  $R$ , the algorithm to compute all the candidate keys is as follows:

$T := \emptyset$

*Main:*

$X := S$  where  $S$  is a super key which does not contain any candidate key in  $T$

remove := true

While remove do

    For each attribute  $A \in X$

        Compute  $\{X-A\}^+$  with respect to  $F$

        If  $\{X-A\}^+$  contains all attributes of  $R$  then

$X := X - \{A\}$

        Else

            remove := false

$T := T \cup X$

Repeat *Main* until no available  $S$  can be found. Finally,  $T$  contains all the candidate keys.

# Normal Forms

Normal Forms for relational databases:

- 1NF, 2NF, 3NF (Codd 1972)
- Boyce-Codd NF (1974)

# First Normal Form (1NF)

This simply means that attribute values are *atomic and* is part of the definition of the relational model.

Atomic: multivalued attributes, composite attributes, and their combinations are disallowed.

There is currently a lot of interests in non-first normal form databases, particularly those where an attribute value can be a table (nested relations).

# Second Normal Form (2NF)

A *prime* attribute is one that is part of a candidate key. Other attributes are *non-prime*.

**Definition:** In an FD  $X \rightarrow Y$ ,  $Y$  is *fully functionally dependent* on  $X$  if there is no  $Z \subset X$  such that  $Z \rightarrow Y$ . Otherwise,  $Y$  is *partially dependent* on  $X$ .

**Definition (Second Normal Form):** A relation scheme is in second normal form (2NF) if **all non-prime attributes are fully functionally dependent** on the relation keys.

A database scheme is in 2NF if all its relations are in 2NF.



# Third Normal Form (3NF)

**Definition** (*Third Normal Form*): A relation scheme is in *third normal form (3NF)* if for all non-trivial FD's of the form  $X \rightarrow A$  that hold, **either  $X$  is a superkey or  $A$  is a prime attribute**.

Note: a FD  $X \rightarrow Y$  is trivial iff  $Y$  is a subset of  $X$ .

*Alternative definition*: A relation scheme is in third normal form if every non-prime attribute is fully functionally dependent on the keys and not transitively dependent on any key.

A database scheme is in 3NF if all its relations are in 3NF.

# Boyce-codd Normal Form (BCNF)

**Definition** (*Boyce-codd Normal Form*):

A relation scheme is in *Boyce-codd Normal Form* (BCNF) if whenever  $X \rightarrow A$  holds (and  $X \rightarrow A$  is non-trivial),  **$X$  is a superkey**.

A database scheme is in BCNF if all its relations are in BCNF.

# On Relational Database Design

1. Anomalies can be removed from relation designs by decomposing them until they are in a normal form.
2. A **decomposition** of a relation scheme,  $R$ , is a set of relation schemes  $\{R_1, \dots, R_n\}$  such that  $R_i \subseteq R$  for each  $i$ , and  $\bigcup_{i=1}^n R_i = R$
3. In a decomposition  $\{R_1, \dots, R_n\}$ , the intersect of each pair of  $R_i$  and  $R_j$  does not have to be empty.
  - Example:  $R = \{A, B, C, D, E\}$ ,  $R_1 = \{A, B\}$ ,  $R_2 = \{A, C\}$ ,  $R_3 = \{C, D, E\}$
4. A naive decomposition: each relation has only attribute.

# Dependency Preserving

A decomposition  $D=\{R_1, \dots, R_n\}$  of  $R$  is **dependency-preserving** wrt a set  $F$  of FDs if:

$$(F_1 \cup \dots \cup F_n)^+ = F^+,$$

where  $F_i$  means the **projection** of  $F$  onto  $R_i$ .

# Lossless-join Decomposition

**A good decomposition should have the following property.**

A decomposition  $\{R_1, \dots, R_n\}$  of  $R$  is a *lossless join* decomposition with respect to a set  $F$  of FD's if for every relation instance  $r$  that satisfies  $F$ :

$$r = \pi_{R_1}(r) \bowtie \dots \bowtie \pi_{R_n}(r).$$

If  $r \subset \pi_{R_1}(r) \bowtie \dots \bowtie \pi_{R_n}(r)$ , the decomposition is *lossy*.

# Test Lossless Join property

This previous test works on **binary** decompositions, below is the general solution to testing lossless join property

Algorithm TEST\_LJ:

1. Create a **matrix**  $S$ , each element  $s_{i,j} \in S$  corresponds the relation  $R_i$  and the attribute  $A_j$ , such that:  $s_{j,i} = a$  if  $A_i \in R_j$ , otherwise  $s_{j,i} = b$ .
2. Repeat the following process until (1)  $S$  has no change OR (2) one row is made up entirely of “a” symbols.
  - i. For each  $X \rightarrow Y$ , choose the rows where the elements corresponding to  $X$  take the value  $a$ .
  - ii. In those chosen rows (must be at least two rows), the elements corresponding to  $Y$  also take the value  $a$  if one of the chosen rows take the value  $a$  on  $Y$ .

Verdict: Decomposition is *lossless* if one row is entirely made up by “a” values.

# Lossless Decomposition into BCNF

## Algorithm TO\_BCNF

$D := \{R_1, R_2, \dots, R_n\}$

**While**  $\exists$  a  $R_i \in D$  and  $R_i$  is not in BCNF **Do**

- { find a  $X \rightarrow Y$  in  $R_i$  that violates BCNF; replace  $R_i$  in  $D$  by  $(R_i - Y)$  and  $(X \cup Y)$ ; }

# Computing a Minimum cover

A set  $F$  of FD's is **minimal** if

1. Every FD  $X \rightarrow Y$  in  $F$  is simple:  $Y$  consists of a single attribute,
2. Every FD  $X \rightarrow A$  in  $F$  is *left-reduced*: there is no proper subset  $Y \subset X$  such that  $X \rightarrow A$  can be replaced with  $Y \rightarrow A$ .

that is, there is no  $Y \subset X$  such that

$$((F - \{X \rightarrow A\}) \cup \{Y \rightarrow A\})^+ = F^+$$

→ i.e., iff  $F \models Y \rightarrow A$

3. No FD in  $F$  can be removed; that is, there is no FD  $X \rightarrow A$  in  $F$  such that

$$(F - \{X \rightarrow A\})^+ = F^+$$

→ i.e., iff  $X \rightarrow A$  is inferred from  $F - \{X \rightarrow A\}$



# Computing a Minimum cover

A *minimal cover* (or *canonical cover*) for  $F$ :

- a minimal set of FD's  $F_{min}$  such that  $F^+ = F_{min}^+$ .

## Algorithm Min\_Cover

Input: a set  $F$  of functional dependencies.

Output: a minimum cover of  $F$ .

Step 1: *Reduce right side.*

Apply Algorithm Reduce\_right to  $F$ .

Step 2: *Reduce left side.*

Apply Algorithm Reduce\_left to the output of Step 2.

Step 3: *Remove redundant FDs.*

Apply Algorithm Remove\_redundancy to the output of Step 2.

The output is a minimum cover.

Next we detail the three algorithms (Reduce\_right, Reduce\_left, Reduce\_redundancy) .

# Computing a Minimum cover (cont)

## 1. Algorithm Reduce\_right

*INPUT:  $F$ .*

*OUTPUT: right side reduced  $F'$ .*

1. For each FD  $X \rightarrow Y \in F$  where  $Y = \{A_1, A_2, \dots, A_k\}$ , we use all  $X \rightarrow \{A_i\}$  (for  $1 \leq i \leq k$ ) to replace  $X \rightarrow Y$ .

## 2. Algorithm Reduce\_left

*INPUT: right side reduced  $F$ .*

*OUTPUT: right and left side reduced  $F'$ .*

1. For each  $X \rightarrow \{A\} \in F$  where  $X = \{A_i : 1 \leq i \leq k\}$ , do the following.  
For  $i = 1$  to  $k$ , replace  $X$  with  $X - \{A_i\}$  if  $A \in (X - \{A_i\})^+$ .

## 3. Algorithm Reduce\_redundancy

*INPUT: right and left side reduced  $F$ .*

*OUTPUT: a minimum cover  $F'$  of  $F$ .*

1. For each FD  $X \rightarrow \{A\} \in F$ , remove it from  $F$  if:  $A \in X^+$  with respect to  $F - \{X \rightarrow \{A\}\}$ .

# Decomposition into 3NF

A **lossless and dependency-preserving decomposition** into 3NF is always possible.

Algorithm 3NF decomposition (Lossless and Dependency-preserving)

1. Find a minimal cover  $G$  for  $F$ .
2. For each left-hand-side  $X$  of a functional dependency that appears in  $G$ , create a relation schema in  $D$  with attributes  $\{X \cup \{A_1\} \cup \{A_2\} \dots \cup \{A_k\}\}$ , where  $X \rightarrow A_1$ ,  $X \rightarrow A_2$ , ...,  $X \rightarrow A_k$  are the only dependencies in  $G$  with  $X$  as left-hand-side ( $X$  is the key to this relation).
3. If none of the relation schemas in  $D$  contains a key of  $R$ , then create one more relation schema in  $D$  that contains attributes that form a key of  $R$ .
4. Eliminate redundant relations from the resulting set of relations in the relational database schema. A relation  $R$  is considered redundant if  $R$  is a projection of another relation  $S$  in the schema; alternately,  $R$  is subsumed by  $S$ .

# Overview: DBMS


Disk, Buffer Replacement Policy

Transaction Management

- ACID properties
- Various schedules: Serializable, Conflict-Serializable, Schedule Graph, Wait for Graph, ...
- Concurrency control (locking)

# Record Format

- Format
  - Fixed-length
  - Variable Length
- Fragmented free space



# Buffer Replacement Policies

## 1. Least Recently Used (LRU)

- release the frame that has not been used for the longest period.
- intuitively appealing idea but can perform badly

## 2. First in First Out (FIFO)

## 3. Most Recently Used (MRU):

- release the frame used most recently

No one is guaranteed to be better than the others. Quite dependent on applications.

# Desirable Properties of Transaction Processing

## ACID Properties

### Atomicity:

A transaction is either performed in its entirety or not performed at all.

### Consistency:

A correct execution of the transaction must take the database from one consistent state to another.

### Isolation:

A transaction should not make its updates visible to other transactions until it is committed.

### Durability/ Permanency:

Once a transaction changes the database and the changes are committed, these changes must never be lost because of subsequent failure.

# Check Conflict Serializability

## *Algorithm*

Step 1: Construct a *schedule* (or ***precedence***) graph – a *directed graph*.

Step 2: Check if the graph is *cyclic*:

- Cyclic: non-serializable.
- Acyclic: serializable.



# Construct a Precedence Graph

Schedule Graph  $GS = (V, A)$  for a schedule  $S$

A vertex in  $V$  represents a transaction.

For two vertices  $T_i$  and  $T_j$ , an arc  $T_i \rightarrow T_j$  is added to  $A$  if

- there are two *conflicting* operations  $O_1 \in T_i$  and  $O_2 \in T_j$ ,
- in  $S$ ,  $O_1$  is before  $O_2$ .

Recall: two operations  $O_1$  and  $O_2$  are *conflicting* if

- they are in different transactions but on the same data item,
- one of them must be a write.

# Locking Rules

In this schema, every transaction  $T$  must obey the following rules.

- 1) If  $T$  has **only one** operation (read/write) manipulating an item  $X$ :
  - obtain a read lock on  $X$  before reading it,
  - obtain a write lock on  $X$  before writing it,
  - unlock  $X$  when done with it.
- 2) If  $T$  has **several** operations manipulating  $X$ :
  - obtain one proper lock only on  $X$ :
    - a read lock if all operations on  $X$  are reads;
    - a write lock if one of these operations on  $X$  is a write.
  - unlock  $X$  after the last operation on  $X$  in  $T$  has been executed.

You should know **Two-Phase Locking!**

Thank you and all the best!